

Hardwarepraktikum Internet-Technologien

Task 8: Software-Defined Networking



Julius-Maximilians-Universität Würzburg

Chair of Computer Science III

A project report submitted by **Group 11**

Enrique Mesonero Ronco

Manuel Calvo Martín

Pablo E. Ortega Ureña

Summer Term 2022

July, 2022

Contents

8. Software-Defined Networking	3
8.2. Installation of the SDN Switch	3
8.3. Simple OpenFlow rules	4
8.4. Replication of a static NAT router with firewall	8

8. Software-Defined Networking

8.2. Installation of the SDN Switch

Our connections were as follows:

SDN Switch

- Port 1 (Internet/Management) to Port 2 of the Switch
- Port 2 to Port 4 of the Switch
- Port 3 to PC A

Switch

- Port 2 to Port 1 (Internet/Management) of the SDN Switch
- Port 3 to PC B (acting as the controller)
- Port 4 to Port 2 of the SDN Switch
- Port 5 to Raspberry Pi

The connection between the controller PC and the SDN Switch is successful. We check that, as it has been just turned on, there are no existing table entries or flows, using the dump-flows command. We also check the manual for understanding future commands.



Figure 8.1: Physical topology of the session

```
root@OpenWrt: ~  
root@OpenWrt:~# ovs-ofctl -h  
ovs-ofctl: OpenFlow switch management utility  
usage: ovs-ofctl [OPTIONS] COMMAND [ARG...]  
  
For OpenFlow switches:  
  show SWITCH                show OpenFlow information  
  dump-desc SWITCH           print switch description  
  dump-tables SWITCH         print table stats  
  dump-table-features SWITCH print table features  
  dump-table-desc SWITCH     print table description (OF1.4+)  
  mod-port SWITCH IFACE ACT  modify port behavior  
  mod-table SWITCH MOD       modify flow table behavior  
                             OF1.1/1.2 MOD: controller, continue, drop  
                             OF1.4+ MOD: evict, noevict, vacancy:low,high, novacancy  
  get-frags SWITCH           print fragment handling behavior  
  set-frags SWITCH FRAG_MODE set fragment handling behavior  
                             FRAG_MODE: normal, drop, reassemble, nx-match  
  dump-ports SWITCH [PORT]   print port statistics  
  dump-ports-desc SWITCH [PORT] print port descriptions  
  dump-flows SWITCH          print all flow entries  
  dump-flows SWITCH FLOW     print matching FLOWS  
  dump-aggregate SWITCH      print aggregate flow statistics  
  dump-aggregate SWITCH FLOW print aggregate stats for FLOWS  
  queue-stats SWITCH [PORT [QUEUE]] dump queue stats  
  add-flow SWITCH FLOW       add flow described by FLOW  
  add-flows SWITCH FILE      add flows from FILE  
  mod-flows SWITCH FLOW      modify actions of matching FLOWS  
  del-flows SWITCH [FLOW]    delete matching FLOWS  
  replace-flows SWITCH FILE  replace flows with those in FILE  
  diff-flows SOURCE1 SOURCE2 compare flows from two sources  
  packet-out SWITCH IN_PORT ACTIONS PACKET...  
                             execute ACTIONS on PACKET  
  monitor SWITCH [MISSLEN] [invalid_ttl] [watch:[...]]  
                             print packets received from SWITCH  
  snoop SWITCH               snoop on SWITCH and its controller  
  add-group SWITCH GROUP     add group described by GROUP  
  add-groups SWITCH FILE     add group from FILE  
  [--may-create] mod-group SWITCH GROUP modify specific group  
  del-groups SWITCH [GROUP]  delete matching GROUPS  
  insert-buckets SWITCH [GROUP] add buckets to GROUP  
  remove-buckets SWITCH [GROUP] remove buckets from GROUP  
  dump-group-features SWITCH print group features
```

Figure 8.2: Checking the manual

```
root@OpenWrt:~# ovs-ofctl dump-flows ovs-br  
root@OpenWrt:~#
```

Figure 8.3: Checking active flows

8.3. Simple OpenFlow rules

We started pingging the Raspberry, the Switch and the PC B from PC A, connected to the SDN. We checked there was no connection yet. After that, we added two flow rules to our SDN, one for each direction of the network. After adding them, the ping was now possible.

```

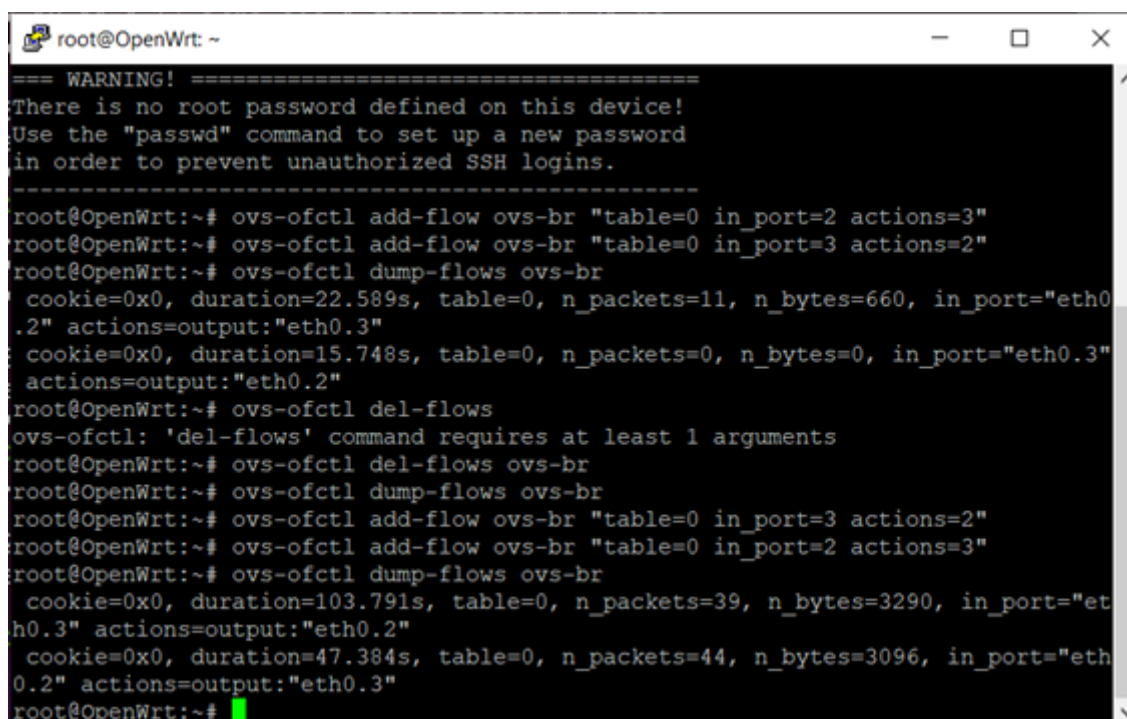
athenyx@athenyx-boreas:~$ ping 10.11.1.1
PING 10.11.1.1 (10.11.1.1) 56(84) bytes of data.
^C
--- 10.11.1.1 ping statistics ---
3 packets transmitted, 0 received, 100% packet loss, time 2053ms

athenyx@athenyx-boreas:~$ ping 10.11.1.3
PING 10.11.1.3 (10.11.1.3) 56(84) bytes of data.
^C
--- 10.11.1.3 ping statistics ---
6 packets transmitted, 0 received, 100% packet loss, time 5110ms

athenyx@athenyx-boreas:~$ ping 10.11.1.5
PING 10.11.1.5 (10.11.1.5) 56(84) bytes of data.
^C
--- 10.11.1.5 ping statistics ---
4 packets transmitted, 0 received, 100% packet loss, time 3054ms

```

Figure 8.4: Pinging the devices from PC A before adding the rules



```

root@OpenWrt: ~
=== WARNING! ===
There is no root password defined on this device!
Use the "passwd" command to set up a new password
in order to prevent unauthorized SSH logins.
=====
root@OpenWrt:~# ovs-ofctl add-flow ovs-br "table=0 in_port=2 actions=3"
root@OpenWrt:~# ovs-ofctl add-flow ovs-br "table=0 in_port=3 actions=2"
root@OpenWrt:~# ovs-ofctl dump-flows ovs-br
  cookie=0x0, duration=22.589s, table=0, n_packets=11, n_bytes=660, in_port="eth0.2" actions=output:"eth0.3"
  cookie=0x0, duration=15.748s, table=0, n_packets=0, n_bytes=0, in_port="eth0.3" actions=output:"eth0.2"
root@OpenWrt:~# ovs-ofctl del-flows
ovs-ofctl: 'del-flows' command requires at least 1 arguments
root@OpenWrt:~# ovs-ofctl del-flows ovs-br
root@OpenWrt:~# ovs-ofctl dump-flows ovs-br
  cookie=0x0, duration=103.791s, table=0, n_packets=39, n_bytes=3290, in_port="eth0.3" actions=output:"eth0.2"
  cookie=0x0, duration=47.384s, table=0, n_packets=44, n_bytes=3096, in_port="eth0.2" actions=output:"eth0.3"
root@OpenWrt:~#

```

Figure 8.5: Rules added, testing the packet flow

We checked that both the packet and the bytes number under the rules increased with each ping and reply.


```

athenyx@athenyx-boreas:~$ ping 10.11.1.1
PING 10.11.1.1 (10.11.1.1) 56(84) bytes of data.
64 bytes from 10.11.1.1: icmp_seq=1 ttl=64 time=3.47 ms
64 bytes from 10.11.1.1: icmp_seq=2 ttl=64 time=0.409 ms
64 bytes from 10.11.1.1: icmp_seq=3 ttl=64 time=0.438 ms
64 bytes from 10.11.1.1: icmp_seq=4 ttl=64 time=0.420 ms
^C
--- 10.11.1.1 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3037ms
rtt min/avg/max/mdev = 0.409/1.185/3.473/1.321 ms
athenyx@athenyx-boreas:~$ ping 10.11.1.3
PING 10.11.1.3 (10.11.1.3) 56(84) bytes of data.
64 bytes from 10.11.1.3: icmp_seq=1 ttl=64 time=3.41 ms
64 bytes from 10.11.1.3: icmp_seq=2 ttl=64 time=0.344 ms
64 bytes from 10.11.1.3: icmp_seq=3 ttl=64 time=0.459 ms
^C
--- 10.11.1.3 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2005ms
rtt min/avg/max/mdev = 0.344/1.404/3.410/1.418 ms
athenyx@athenyx-boreas:~$ ping 10.11.1.5
PING 10.11.1.5 (10.11.1.5) 56(84) bytes of data.
64 bytes from 10.11.1.5: icmp_seq=1 ttl=64 time=6.49 ms
64 bytes from 10.11.1.5: icmp_seq=2 ttl=64 time=2.46 ms
64 bytes from 10.11.1.5: icmp_seq=3 ttl=64 time=2.40 ms
^C
--- 10.11.1.5 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2003ms
rtt min/avg/max/mdev = 2.403/3.783/6.490/1.914 ms

```

Figure 8.6: Pinging the devices from PC A after adding the rules

Following the next step, we deleted the existing flows and added new ones based on the MAC addresses of our devices. Since ARP requests should be also transmitted correctly, we added an additional rule where packets with MAC address equal to the broadcast address must be flooded.

```

root@OpenWrt:~# ovs-ofctl del-flows ovs-br
root@OpenWrt:~# ovs-ofctl add-flow ovs-br "table=0 dl_dst=e4:5f:01:05:f9:e0 action=2"
root@OpenWrt:~# ovs-ofctl add-flow ovs-br "table=0 dl_dst=0c:9d:92:c7:4f:e5 action=3"
root@OpenWrt:~# ovs-ofctl add-flow ovs-br "table=0 dl_dst=08:00:27:6f:0c:ef action=2"
root@OpenWrt:~# ovs-ofctl add-flow ovs-br "table=0 dl_dst=ff:ff:ff:ff:ff:ff action=flood"
root@OpenWrt:~# ovs-ofctl add-flow ovs-br "table=0 dl_dst=64:d1:54:9a:4a:71 action=2"

```

Figure 8.7: New layer 2 flow rules added

The connection was again successfully established without issues.

Afterwards, we removed those rules and tried forwarding based on layer 3. We retained the flood on broadcast rule, but changed those referring to MAC addresses to their IP address.

```

athenyx@athenyx-boreas:~$ ping 10.11.1.5
PING 10.11.1.5 (10.11.1.5) 56(84) bytes of data.
From 10.11.1.4 icmp_seq=1 Destination Host Unreachable
From 10.11.1.4 icmp_seq=2 Destination Host Unreachable
From 10.11.1.4 icmp_seq=3 Destination Host Unreachable
From 10.11.1.4 icmp_seq=4 Destination Host Unreachable
From 10.11.1.4 icmp_seq=5 Destination Host Unreachable
From 10.11.1.4 icmp_seq=6 Destination Host Unreachable
From 10.11.1.4 icmp_seq=8 Destination Host Unreachable
From 10.11.1.4 icmp_seq=9 Destination Host Unreachable
64 bytes from 10.11.1.5: icmp_seq=17 ttl=64 time=2049 ms
64 bytes from 10.11.1.5: icmp_seq=18 ttl=64 time=1025 ms
From 10.11.1.4 icmp_seq=19 Destination Host Unreachable
From 10.11.1.4 icmp_seq=20 Destination Host Unreachable
From 10.11.1.4 icmp_seq=21 Destination Host Unreachable
From 10.11.1.4 icmp_seq=22 Destination Host Unreachable
From 10.11.1.4 icmp_seq=23 Destination Host Unreachable
From 10.11.1.4 icmp_seq=24 Destination Host Unreachable
From 10.11.1.4 icmp_seq=25 Destination Host Unreachable
^C
--- 10.11.1.5 ping statistics ---
35 packets transmitted, 2 received, +15 errors, 94.2857% packet loss, time 34728 ms
rtt min/avg/max/mdev = 1025.226/1537.225/2049.225/511.999 ms, pipe 4

```

Figure 8.8: Awkward behavior of the ping after the IP flow rules

```

root@OpenWrt:~# ovs-ofctl dump-flows ovs-br
cookie=0x0, duration=707.232s, table=0, n_packets=0, n_bytes=0, ip,nw_dst=10.11.1.3 actions=output:"eth0.2",dec_ttl
root@OpenWrt:~# ovs-ofctl dump-flows ovs-br
cookie=0x0, duration=736.268s, table=0, n_packets=216, n_bytes=17160, dl_dst=ff:ff:ff:ff:ff:ff actions=FLOOD
cookie=0x0, duration=730.659s, table=0, n_packets=22, n_bytes=2156, ip,nw_dst=10.11.1.1 actions=output:"eth0.2",dec_ttl
cookie=0x0, duration=725.713s, table=0, n_packets=85, n_bytes=8330, ip,nw_dst=10.11.1.4 actions=output:"eth0.3",dec_ttl
cookie=0x0, duration=720.466s, table=0, n_packets=0, n_bytes=0, ip,nw_dst=10.11.1.5 actions=output:"eth0.2",dec_ttl
cookie=0x0, duration=715.756s, table=0, n_packets=0, n_bytes=0, ip,nw_dst=10.11.1.3 actions=output:"eth0.2",dec_ttl
root@OpenWrt:~# ovs-ofctl dump-flows ovs-br
cookie=0x0, duration=741.322s, table=0, n_packets=223, n_bytes=18144, dl_dst=ff:ff:ff:ff:ff:ff actions=FLOOD
cookie=0x0, duration=735.714s, table=0, n_packets=31, n_bytes=3038, ip,nw_dst=10.11.1.1 actions=output:"eth0.2",dec_ttl
cookie=0x0, duration=730.768s, table=0, n_packets=94, n_bytes=9212, ip,nw_dst=10.11.1.4 actions=output:"eth0.3",dec_ttl
cookie=0x0, duration=725.521s, table=0, n_packets=0, n_bytes=0, ip,nw_dst=10.11.1.5 actions=output:"eth0.2",dec_ttl
cookie=0x0, duration=720.811s, table=0, n_packets=0, n_bytes=0, ip,nw_dst=10.11.1.3 actions=output:"eth0.2",dec_ttl

```

Figure 8.9: IP flow rules and their data flow

This configuration failed most of the time, but looking at Wireshark we realized our problem. The ARP request was reaching the computer which we were trying to obtain the MAC address from, thanks to the flood on broadcast rule. However, since the ARP reply was directed to a specific MAC address, once it reached the SDN switch, there were no rules regarding that MAC address, only the IP address, so the ARP reply got discarded.

It was then, after looking at the OVS fields document, that we found out that our SDN could look at the packet, and if it was ARP, check the target IP address contained in that packet.

```

ovs-ofctl add-flow ovs-br "table=0 ip nw_dst=10.11.1.1 priority=2 action=dec_ttl,2" #
ovs-ofctl add-flow ovs-br "table=0 ip nw_dst=10.11.1.4 priority=2 action=dec_ttl,3" #
ovs-ofctl add-flow ovs-br "table=0 ip nw_dst=10.11.1.5 priority=2 action=dec_ttl,2" #
ovs-ofctl add-flow ovs-br "table=0 ip nw_dst=10.11.1.3 priority=2 action=dec_ttl,2" #
ovs-ofctl add-flow ovs-br "table=0 ip nw_dst=10.11.1.2 priority=2 action=dec_ttl,2" #

```

Figure 8.10: IP flow rules

```

ovs-ofctl add-flow ovs-br "table=0 arp arp_tpa=10.11.1.1 priority=1 action=2" #
ovs-ofctl add-flow ovs-br "table=0 arp arp_tpa=10.11.1.4 priority=1 action=3" #
ovs-ofctl add-flow ovs-br "table=0 arp arp_tpa=10.11.1.5 priority=1 action=2" #
ovs-ofctl add-flow ovs-br "table=0 arp arp_tpa=10.11.1.3 priority=1 action=2" #
ovs-ofctl add-flow ovs-br "table=0 arp arp_tpa=10.11.1.2 priority=1 action=2" #

```

Figure 8.11: ARP flow rules

With these settings, our pings started working again. We were also able to check that if we set the TTL of the packet under 2, the packet would get discarded by the SDN and never reach its target for a response.

576	2453.7199915...	e4:5f:01:05:f9:e0	Tp-LinkT_d0:25:ac	ARP	60 Who has 10.11.1.4? Tell 10.11.1.1
577	2453.7200116...	Tp-LinkT_d0:25:ac	e4:5f:01:05:f9:e0	ARP	42 10.11.1.4 is at 50:3e:aa:d0:25:ac
578	2454.5945553...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0025, seq=1/256, ttl=1 (no response...
579	2455.5957198...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0025, seq=2/512, ttl=1 (no response...
580	2456.6196914...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0025, seq=3/768, ttl=1 (no response...
581	2457.6437436...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0025, seq=4/1024, ttl=1 (no respons...
582	2458.6677248...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0025, seq=5/1280, ttl=1 (no respons...
583	2459.6916835...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0025, seq=6/1536, ttl=1 (no respons...
584	2460.7157128...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0025, seq=7/1792, ttl=1 (no respons...
585	2461.7396750...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0025, seq=8/2048, ttl=1 (no respons...
586	2462.7637085...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0025, seq=9/2304, ttl=1 (no respons...
587	2463.7876993...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0025, seq=10/2560, ttl=1 (no respon...
588	2464.8116840...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0025, seq=11/2816, ttl=1 (no respon...
589	2468.8184743...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0026, seq=1/256, ttl=2 (reply in 59...
590	2468.8208133...	10.11.1.1	10.11.1.4	ICMP	98 Echo (ping) reply id=0x0026, seq=1/256, ttl=63 (request in...
591	2469.8201745...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0026, seq=2/512, ttl=2 (reply in 59...
592	2469.8221845...	10.11.1.1	10.11.1.4	ICMP	98 Echo (ping) reply id=0x0026, seq=2/512, ttl=63 (request in...
593	2470.8214266...	10.11.1.4	10.11.1.1	ICMP	98 Echo (ping) request id=0x0026, seq=3/768, ttl=2 (reply in 59...
594	2470.8230950...	10.11.1.1	10.11.1.4	ICMP	98 Echo (ping) reply id=0x0026, seq=3/768, ttl=63 (request in...

Figure 8.11: Replies to ICMP in Wireshark with different TTLs

8.4. Replication of a static NAT router with firewall

For this subtask, we adjusted the IP address of PC A from 10.11.1.4 to 10.11.123.1. We also set 10.11.123.254 as its gateway, and also changed the Raspberry Pi gateway to be 10.11.1.254.

```

Verbindungsgeschwindigkeit 1000 Mb/s
IPv4-Adresse 10.11.123.1
IPv6-Adresse fe80::cff8:a3f1:b007:4c72
Geräteadresse 50:3E:AA:D0:25:AC
Vorgabestrecke 10.11.123.254
DNS

```

Figure 8.12: New IP and gateway for PC A

The connection didn't work yet between the two subnets, as there is no way for the ARP requests to reach the other port.

```
ovs-ofctl add-flow ovs-br "table=0 in_port=2 priority=100 arp nw_dst=10.11.1.254
actions=load:0x2->arp_op,move:arp_spa->arp_tpa,move:arp_sha->arp_tha,move:eth_src->eth_dst,
set_field:64:d1:54:22:22:22->eth_src,set_field:64:d1:54:22:22:22->arp_sha,set_field:10.
11.1.254->arp_spa,output:in_port" #
ovs-ofctl add-flow ovs-br "table=0 in_port=3 priority=100 arp nw_dst=10.11.123.254
actions=load:0x2->arp_op,move:arp_spa->arp_tpa,move:arp_sha->arp_tha,move:eth_src->eth_dst,
set_field:64:d1:54:33:33:33->eth_src,set_field:64:d1:54:33:33:33->arp_sha,set_field:10.
11.123.254->arp_spa,output:in_port" #
```

Figure 8.13: ARP configuration

For that communication to be made, we used the two flow rules above, where whenever a connection from the switch port (Port 2) was directed to the gateway IP, the SDN would set-up an ARP reply (as indicated by the 2 opcode), directed to the MAC address of the sender (changed to target now), with the made-up address of our port 2, and the gateway address. The SDN sends that through the same port. This way, now every device connected to the switch knows about the SDN switch gateway. In similar fashion, another rule is made for port 3, but with the subnet 10.11.123.0/24.

Now that the devices can reach the gateways, the SDN router only needs to know (as this is proactive forwarding) what to do whenever a packet from our end devices reaches those gateways. We used PC A and the Raspberry for this.

```
ovs-ofctl add-flow ovs-br "table=0 ip nw_dst=10.11.1.0/24
action=set_field:64:d1:54:22:22:22->eth_src,set_field:e4:5f:01:05:f9:e0->eth_dst,
dec_ttl,2" #
ovs-ofctl add-flow ovs-br "table=0 ip nw_dst=10.11.123.0/24
action=set_field:64:d1:54:33:33:33->eth_src,set_field:50:3e:aa:d0:25:ac->eth_dst,
dec_ttl,3" #
```

Figure 8.14: IP configuration

These are pretty similar to the IP flow rules of the previous subtask. However now, the layer 2 source and destination addresses are set, with the MAC addresses of the used port and the end device, respectively.

53	106.059655115	Tp-LinkT_d0:25:ac	Broadcast	ARP	42 Who has 10.11.123.254? Tell 10.11.123.1
54	106.062173093	Routerbo_33:33:33	Tp-LinkT_d0:25:ac	ARP	60 10.11.123.254 is at 64:d1:54:33:33:33
55	106.062193256	10.11.123.1	10.11.1.1	ICMP	98 Echo (ping) request id=0x002a, seq=1/256, ttl=64 (reply in 5
56	106.064540368	10.11.1.1	10.11.123.1	ICMP	98 Echo (ping) reply id=0x002a, seq=1/256, ttl=63 (request in
57	107.060780195	10.11.123.1	10.11.1.1	ICMP	98 Echo (ping) request id=0x002a, seq=2/512, ttl=64 (reply in 5
58	107.063050774	10.11.1.1	10.11.123.1	ICMP	98 Echo (ping) reply id=0x002a, seq=2/512, ttl=63 (request in
59	108.062330759	10.11.123.1	10.11.1.1	ICMP	98 Echo (ping) request id=0x002a, seq=3/768, ttl=64 (reply in 6
60	108.063020960	10.11.1.1	10.11.123.1	ICMP	98 Echo (ping) reply id=0x002a, seq=3/768, ttl=63 (request in
61	109.071243838	10.11.123.1	10.11.1.1	ICMP	98 Echo (ping) request id=0x002a, seq=4/1024, ttl=64 (reply in
62	109.073387366	10.11.1.1	10.11.123.1	ICMP	98 Echo (ping) reply id=0x002a, seq=4/1024, ttl=63 (request in
63	110.072621837	10.11.123.1	10.11.1.1	ICMP	98 Echo (ping) request id=0x002a, seq=5/1280, ttl=64 (reply in
64	110.073409792	10.11.1.1	10.11.123.1	ICMP	98 Echo (ping) reply id=0x002a, seq=5/1280, ttl=63 (request in
65	111.083234095	10.11.123.1	10.11.1.1	ICMP	98 Echo (ping) request id=0x002a, seq=6/1536, ttl=64 (reply in
66	111.085595758	10.11.1.1	10.11.123.1	ICMP	98 Echo (ping) reply id=0x002a, seq=6/1536, ttl=63 (request in
67	112.084866347	10.11.123.1	10.11.1.1	ICMP	98 Echo (ping) request id=0x002a, seq=7/1792, ttl=64 (reply in
68	112.087192351	10.11.1.1	10.11.123.1	ICMP	98 Echo (ping) reply id=0x002a, seq=7/1792, ttl=63 (request in
69	129.327078878	fe80::cff8:a3f1:b00...	ff02::fb	MDNS	203 Standard query 0x0000 PTR _ipps._tcp.local, "QM" question PTR
70	129.327248854	10.11.123.1	224.0.0.251	MDNS	183 Standard query 0x0000 PTR _ipps._tcp.local, "QM" question PTR
71	129.963995949	10.11.123.1	10.11.1.1	ICMP	98 Echo (ping) request id=0x002b, seq=1/256, ttl=1 (no response
72	130.987286703	10.11.123.1	10.11.1.1	ICMP	98 Echo (ping) request id=0x002b, seq=2/512, ttl=1 (no response
73	132.011260266	10.11.123.1	10.11.1.1	ICMP	98 Echo (ping) request id=0x002b, seq=3/768, ttl=1 (no response
74	133.035287846	10.11.123.1	10.11.1.1	ICMP	98 Echo (ping) request id=0x002b, seq=4/1024, ttl=1 (no response

Figure 8.15: Checking with Wireshark

With this, communication is again restored. We also checked that with a TTL of 1, it still wouldn't work. This matches what our router did in previous sessions, and what ofproto/trace tells us.

```

root@OpenWrt: ~
ovs-appctl: ovs-vswitchd: server returned an error
root@OpenWrt:~# ovs-appctl ofproto/trace ovs-br in_port=3,icmp,nw_src=10.11.123.1,nw_dst=10.11.1.1
Flow: icmp,in_port=3,vlan_tci=0x0000,dl_src=00:00:00:00:00:00,dl_dst=00:00:00:00:00:00,nw_src=10.11.123.1,nw_dst=10.11.1.1,nw_tos=0,nw_ecn=0,nw_ttl=0,icmp_type=0,icmp_code=0

bridge("ovs-br")
-----
0. ip,nw_dst=10.11.1.0/24, priority 32768
   mod_dl_src:64:d1:54:22:22:22
   mod_dl_dst:e4:5f:01:05:f9:e0
   dec_ttl
   >> IPv4 decrement TTL exception

Final flow: icmp,in_port=3,vlan_tci=0x0000,dl_src=64:d1:54:22:22:22,dl_dst=e4:5f:01:05:f9:e0,nw_src=10.11.123.1,nw_dst=10.11.1.1,nw_tos=0,nw_ecn=0,nw_ttl=0,icmp_type=0,icmp_code=0
MegafLOW: recirc_id=0,eth,ip,in_port=3,dl_src=00:00:00:00:00:00,dl_dst=00:00:00:00:00:00,nw_src=8.0.0.0/5,nw_dst=10.11.1.0/25,nw_ttl=0,nw_frag=no
Datapath actions: set(eth(src=64:d1:54:22:22:22,dst=e4:5f:01:05:f9:e0)),userspace(pid=2255733610,controller(reason=2,dont_send=1,continuation=0,recirc_id=80,rule_cookie=0,controller_id=0,max_len=65535))
root@OpenWrt:~#

```

Figure 8.16: Ofproto/trace result

Next, we moved towards the NAT/Firewall part. For that, we set that any packet directed to the subnet 10.11.1.0/24, incoming from port 3 (PC A's port) would be drop, unless of course, it was directed specifically to the gateway's address (10.11.1.254) and to port 12221. Whenever the NAT would receive this kind of packet, it would change the IP field to the Raspberry one, and the port to 22, the SSH one.

This must be bidirectional, so we also set that any packet incoming from port 2 towards PC A, would have its IP masked to the gateway one, and the port changed back to 12221. It at first didn't work for us, as we forgot to add this

to the previous rules and instead made new rules. Since both rules matched, but were different, and didn't have any priority attached, only the first added rule (which was more general) was acting. However, to solve this, we combined both the previous IP flow rules and the NAT/Firewall rules. The SSH to port 12221 was now working, with all other connections inside the 10.11.1.0/24 subnet rejected, and the IP correctly masked behind the NAT.

```
ovs-ofctl add-flow ovs-br "table=0 in_port=3 tcp nw_dst=10.11.1.254 tcp_dst=12221
action=set_field:10.11.1.1->nw_dst,set_field:22->tcp_dst,set_field:64:d1:54:22:22->eth
_src,set_field:e4:5f:01:05:f9:e0->eth_dst,dec_ttl,2" #

ovs-ofctl add-flow ovs-br "table=0 in_port=3 tcp nw_dst=10.11.1.0/24 action=drop" #
ovs-ofctl add-flow ovs-br "table=0 in_port=2 tcp nw_dst=10.11.123.0/24
action=set_field:64:d1:54:33:33:33->eth_src,set_field:50:3e:aa:d0:25:ac->eth_dst,set_fiel
d:12221->tcp_src,dec_ttl,set_field:10.11.1.254->nw_src,3" #
```

Figure 8.17: New TCP rules for NAT

51	172.452880958	10.11.123.1	10.11.1.254	TCP	74	59780	→	12221	[SYN]	Seq=0	Win=64240	Len=0	MSS=1460	SACK_PERM=...
52	172.456325021	10.11.1.254	10.11.123.1	TCP	74	12221	→	59780	[SYN, ACK]	Seq=0	Ack=1	Win=65160	Len=0	MSS=1460...
53	172.456341322	10.11.123.1	10.11.1.254	TCP	66	59780	→	12221	[ACK]	Seq=1	Ack=1	Win=64256	Len=0	TSval=1013073...
54	172.456528203	10.11.123.1	10.11.1.254	TCP	107	59780	→	12221	[PSH, ACK]	Seq=1	Ack=1	Win=64256	Len=41	TSval=1...
55	172.456985628	10.11.1.254	10.11.123.1	TCP	66	12221	→	59780	[ACK]	Seq=1	Ack=42	Win=65152	Len=0	TSval=547591...
56	172.509468915	10.11.1.254	10.11.123.1	TCP	114	12221	→	59780	[PSH, ACK]	Seq=1	Ack=42	Win=65152	Len=48	TSval=...
57	172.509497392	10.11.123.1	10.11.1.254	TCP	66	59780	→	12221	[ACK]	Seq=42	Ack=49	Win=64256	Len=0	TSval=10130...
58	172.510029277	10.11.123.1	10.11.1.254	TCP	1578	59780	→	12221	[PSH, ACK]	Seq=42	Ack=49	Win=64256	Len=1512	TSv...

Figure 8.18: TCP connection established with the new port and gateway

```
hwp@hwp-l:~$ ssh pi@10.11.1.254 -p 12221
The authenticity of host '[10.11.1.254]:12221 ([10.11.1.254]:12221)' can't be established.
ECDSA key fingerprint is SHA256:oh80ILNnbcnum7407C2C/kJdQfqlvNm7tpJwG1L6ly0.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '[10.11.1.254]:12221' (ECDSA) to the list of known hosts.
pi@10.11.1.254's password:
Linux raspberrypi 5.10.17-v7l+ #1414 SMP Fri Apr 30 13:20:47 BST 2021 armv7l

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by applicable law.
Last login: Wed May 18 13:44:43 2022 from 10.11.123.1
pi@raspberrypi:~$ ^C
pi@raspberrypi:~$ exit
Abgemeldet
Connection to 10.11.1.254 closed.
hwp@hwp-l:~$ ping 10.11.1.1
PING 10.11.1.1 (10.11.1.1) 56(84) Bytes Daten.
^C
--- 10.11.1.1 ping statistics ---
4 Pakete übertragen, 0 empfangen, 100% Paketverlust, Zeit 3059ms

hwp@hwp-l:~$ ping 10.11.1.3
PING 10.11.1.3 (10.11.1.3) 56(84) Bytes Daten.
^C
--- 10.11.1.3 ping statistics ---
3 Pakete übertragen, 0 empfangen, 100% Paketverlust, Zeit 2051ms

hwp@hwp-l:~$ ping 10.11.1.5
PING 10.11.1.5 (10.11.1.5) 56(84) Bytes Daten.
^C
--- 10.11.1.5 ping statistics ---
3 Pakete übertragen, 0 empfangen, 100% Paketverlust, Zeit 2027ms
```

Figure 8.19: Pings to different devices and SSH success

Finally, we also made another rule to be able to access PC B from PC A, on port 10810, in order to read the temperature sensor data. At first it didn't work for us, even when using the same configuration used to connect to the Raspberry Pi, and making sure the IP and MAC addresses were correct. The connection kept resetting.

33	84.188661862	10.11.123.1	10.11.1.254	TCP	74 59778 → 12221 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 SACK_PERM=1
34	84.192463811	10.11.1.254	10.11.123.1	TCP	74 22 → 59778 [SYN, ACK] Seq=0 Ack=1 Win=65160 Len=0 MSS=1460 SACK_PERM=1
35	84.192481243	10.11.123.1	10.11.1.254	TCP	54 59778 → 22 [RST] Seq=1 Win=0 Len=0
36	85.194215771	10.11.123.1	10.11.1.254	TCP	74 [TCP Retransmission] 59778 → 12221 [SYN] Seq=0 Win=64240 Len=0
37	85.195816591	10.11.1.254	10.11.123.1	TCP	74 [TCP Retransmission] 22 → 59778 [SYN, ACK] Seq=0 Ack=1 Win=65160
38	85.195853268	10.11.123.1	10.11.1.254	TCP	54 59778 → 22 [RST] Seq=1 Win=0 Len=0
39	86.264368056	10.11.1.254	10.11.123.1	TCP	74 [TCP Retransmission] 22 → 59778 [SYN, ACK] Seq=0 Ack=1 Win=65160
40	86.264411376	10.11.123.1	10.11.1.254	TCP	54 59778 → 22 [RST] Seq=1 Win=0 Len=0
41	88.344331933	10.11.1.254	10.11.123.1	TCP	74 [TCP Retransmission] 22 → 59778 [SYN, ACK] Seq=0 Ack=1 Win=65160
42	88.344375494	10.11.123.1	10.11.1.254	TCP	54 59778 → 22 [RST] Seq=1 Win=0 Len=0
43	89.386212499	Tp-LinkT_d0:25:ac	Routerbo_33:33:33	ARP	42 Who has 10.11.123.254? Tell 10.11.123.1
44	89.388651482	Routerbo_33:33:33	Tp-LinkT_d0:25:ac	ARP	60 10.11.123.254 is at 64:d1:54:33:33:33
45	92.424732420	10.11.1.254	10.11.123.1	TCP	74 [TCP Retransmission] 22 → 59778 [SYN, ACK] Seq=0 Ack=1 Win=65160
46	92.424771154	10.11.123.1	10.11.1.254	TCP	54 59778 → 22 [RST] Seq=1 Win=0 Len=0
47	100.504240978	10.11.1.254	10.11.123.1	TCP	74 [TCP Retransmission] 22 → 59778 [SYN, ACK] Seq=0 Ack=1 Win=65160
48	100.504282570	10.11.123.1	10.11.1.254	TCP	54 59778 → 22 [RST] Seq=1 Win=0 Len=0
49	117.144738561	10.11.1.254	10.11.123.1	TCP	74 [TCP Retransmission] 22 → 59778 [SYN, ACK] Seq=0 Ack=1 Win=65160
50	117.144767745	10.11.123.1	10.11.1.254	TCP	54 59778 → 22 [RST] Seq=1 Win=0 Len=0
51	172.452880958	10.11.123.1	10.11.1.254	TCP	74 59780 → 12221 [SYN] Seq=0 Win=64240 Len=0 MSS=1460 SACK_PERM=1

Figure 8.20: Failure to connect to the DASH server

Figure 8.20 shows the same kind of error we got; however, these were directed to port 22. When it was instead 8050 (DASH server) or 10810, the same errors popped up.

After some help, we fixed it, realizing it was an error coming from the DASH server file, as it didn't permit connections from outside networks. To fix it, we added `host = '0.0.0.0'` to the `app.run_server` line. With that, the connection was successful.

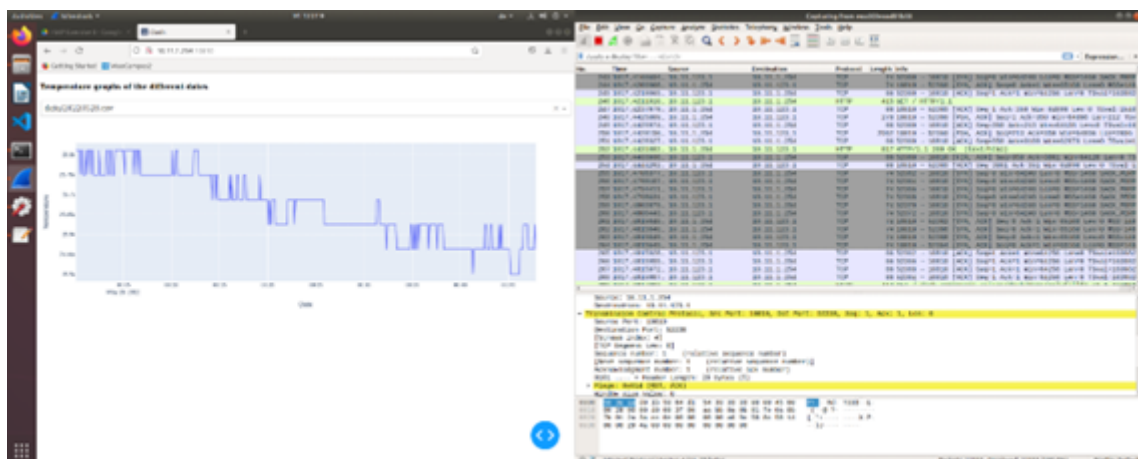


Figure 8.21: Success on connecting to the DASH server